Incremental Support Vector Learning:
Analysis, Implementation and Applications

Pavel Laskov   LASKOV@FIRST.FHG.DE
Christian Gehl   CGEHL@FIRST.FHG.DE
Stefan Krüger   KRUEGERS@FIRST.FHG.DE
Fraunhofer-FIRST.IDA
Kekuléstrasse 7
12489 Berlin, Germany

Klaus-Robert Müller   KLAUS@FIRST.FHG.DE
Fraunhofer-FIRST.IDA
Kekuléstrasse 7
12489 Berlin, Germany
and
University of Potsdam
August-Bebelstrasse 89
14482 Potsdam, Germany

Abstract

Incremental Support Vector Machines (SVM) are instrumental in practical applications of online learning. This work focuses on the design and analysis of efficient incremental SVM learning, with the aim of providing a fast, numerically stable and robust implementation. A detailed analysis of convergence and of algorithmic complexity of incremental SVM learning is carried out. Based on this analysis, a new design of storage and numerical operations is proposed, which speeds up the training of an incremental SVM by a factor of 5 to 20. The performance of the new algorithm is demonstrated in two scenarios: learning with limited resources and active learning. Various applications of the algorithm, such as in drug discovery, online monitoring of industrial devices and and surveillance of network traffic, can be foreseen.

Keywords: incremental SVM, online learning, drug discovery, intrusion detection

1. Introduction

Online learning is a classical learning scenario in which training data is provided one example at a time, as opposed to the batch mode in which all examples are available at once (e.g. Robbins and Munro (1951); Murata (1992); Saad (1998); Bishop (1995); Orr and Müller (1998); LeCun et al. (1998); Murata et al. (2002)).

Online learning is advantageous when dealing with (a) very large or (b) non-stationary data. In the case of non-stationary data, batch algorithms will generally fail if ambiguous information, e.g. different distributions varying over time, is present and is erroneously integrated by the batch algorithm (cf. Murata (1992); Murata et al. (2002)). Many problems of high interest in machine learning can be naturally viewed as online ones. An important practical advantage of online algorithms is

that they allow to incorporate additional training data, when it is available, without re-training from scratch. Given that training is usually the most computationally intensive task, it is not surprising that availability of online algorithms is a major pre-requisite imposed by practitioners that work on large data sets (cf. LeCun et al. (1998)) or even have to perform real-time estimation tasks for continuous data streams, such as in intrusion detection (e.g. Laskov et al. (2004); Eskin et al. (2002)), web-mining (e.g. Chakrabarti (2002)) or brain computer interfacing (e.g. Blankertz et al. (2003)).

In the 1980's online algorithms were investigated in the context of PAC learning (e.g. Angluin (1988); Littlestone et al. (1991)). With the emergence of Support Vector Machines (SVM) in the mid-1990's, interest to online algorithms for this learning method arose as well. However, early work on this subject (e.g. Syed et al., 1999; Rüping, 2002; Kivinen et al., 2001; Ralaivola and d'Alché Buc, 2001) provided only approximate solutions.

An exact solution to the problem of online SVM learning has been found by Cauwenberghs and Poggio (2001). Their incremental algorithm (hereinafter referred to as a C&P algorithm) updates an optimal solution of an SVM training problem after one training example is added (or removed).

Unfortunately acceptance of the C&P algorithm in the machine learning community has been somewhat marginal, not to mention that it remains widely unknown to potential practitioners. Only a handful of follow-up publications is known that extend this algorithm to other related learning problems (e.g. Martin, 2002; Ma et al., 2003; Tax and Laskov, 2003); to our knowledge, no successful practical applications of this algorithm has been reported.

At a first glance, a limited interest to incremental SVM learning may seem to result the absence of well-accepted implementations, such as its counterparts SVM$^{\text{light}}$ (Joachims, 1999), SMO (Platt, 1999) and LIBSVM (Chang and Lin, 2000) for batch SVM learning. The original Matlab implementation by the authors[1] has essentially the semantics of batch learning: training examples are loaded all at once (although learned one at a time) and unlearning is only used for computation of the – ingenious – leave-one-out bound.

There are, however, deeper reasons why incremental SVM may not be so easy to implement. To understand them – and to build a foundation for an efficient design and implementation of the algorithm, a detailed analysis of the incremental SVM technique is carried out in this paper. In particular we address the "accounting" details, which contain pitfalls of performance bottlenecks unless underlying data structures are carefully designed, and analyze convergence of the algorithm.

The following are the main results of our analysis:

1. Computational complexity of a minor iteration of the algorithm is quadratic in the number of training examples learned so far. The actual runtime depends on the balance of memory access and arithmetic operations in a minor iteration.

2. The main incremental step of the algorithm is guaranteed to bring progress in the objective function if a kernel matrix is positive semi-definite.

Based on the results of our analysis, we propose a new storage design and organization of computation for a minor iteration of the algorithm. The idea is to judiciously use row-major and column-major storage of matrices, instead of one-dimensional arrays, in order to possibly eliminate selection operations. The second building block of our design is gaxpy-type matrix-vector multiplication, which allows to further minimize selection operations which cannot be eliminated by storage

---

1. http://bach.ece.jhu.edu/pub/gert/svm/incremental/

design alone. Our experiments show that the new design improves computational efficiency by the factor of 5 to 20.

To demonstrate applicability of incremental SVM to practical applications, two learning scenarios are presented. Learning with limited resources allows to learn from large data sets with as little as 2% of the data needed to be stored in memory. Active learning is another powerful technique by means of which learning can be efficiently carried out in large data sets with limited availability of labels. Various applications of the algorithm, such as in drug discovery, online monitoring of industrial devices and and surveillance of network traffic, can be foreseen.

In order to make this contribution self-contained, we begin with the presentation of the C&P algorithm, highlighting the details that are necessary for a subsequent analysis. An extension of the basic incremental SVM to one-class classification is presented in Section 3. Convergence analysis of the algorithm is carried out in Section 4. Analysis of computational complexity, design of efficient storage and organization of operations are presented and evaluated in Section 5. Finally, potential applications of incremental SVM for learning with limited resources and active learning are illustrated in Section 6.

2. Incremental SVM Algorithm

In this section we present the basic incremental SVM algorithm. Before proceeding with our presentation we need to establish some notation.

2.1 Preliminaries

We assume the training data and their labels are given by a set

$$\{(x_1, y_1), \ldots, (x_n, y_n)\}.$$

The inner product between data points in a feature space is defined by a kernel function $k(x_i, x_j)$. The $n \times n$ kernel matrix K^0 contains the inner product values for all $1 \leq i, j \leq n$. The matrix K is obtained from the kernel matrix by incorporating the labels:

$$K = K^0 \odot (yy^T).$$

The operator $\odot$ denotes the element-wise matrix product, and a vector y denotes labels as an $n \times 1$ vector. Using this notation, the SVM training problem can be formulated as

$$\max_{\mu} \min_{\substack{0 \leq \alpha \leq C \\ y^T \alpha = 0}} W := -1^T \alpha + \frac{C}{2} \alpha^T K \alpha + \mu y^T \alpha. \tag{1}$$

Unlike the classical setting of the SVM training problem, which is usually formulated as maximization or minimization, the problem (1) is a saddle-point formulation obtained by incorporating the equality constraint directly into the cost function. The reason for such construction will become clear shortly.

2.2 Derivation of the Basic Incremental SVM Algorithm

The main building block of the incremental SVM is a procedure for adding one example to an existing optimal solution. When a new point x_c is added, its weight α_c is initially set to 0. If this

assignment is not an optimal solution, i.e. when x_c should become a support vector, the weights of other points and the threshold μ must be updated in order to obtain an optimal solution for the enlarged data set. The procedure can be reversed for a removal of an example: its weight is forced to zero while updating weights of the remaining examples and the threshold μ so that the solution obtained with $\alpha_c = 0$ is optimal for the reduced data set. For the remaining part of this paper we only consider addition of examples.

The saddle point of the problem (1) is given by the Kuhn-Tucker conditions:

$$g_i := -1 + K_{i,:}\alpha + \mu y_i \begin{cases} \geq 0, & \text{if } \alpha_i = 0 \\ = 0, & \text{if } 0 < \alpha_i < C \\ \leq 0, & \text{if } \alpha_i = C \end{cases} \tag{2}$$

$$\frac{\partial W}{\partial \mu} := y^T \alpha = 0. \tag{3}$$

Before an addition of a new example x_c, the Kuhn-Tucker conditions are satisfied for all previous examples. The goal of the weight update in the incremental SVM algorithm is to find a weight assignment such that the Kuhn-Tucker conditions are satisfied for the enlarged data set.

Let us introduce some further notation. Let the set S denote unbounded support vectors ($0 < \alpha_i < C$), the set E denote bounded support vectors ($\alpha_i = C$), and the set O denote non-support vectors ($\alpha_i = 0$); let $R = E \cup O$. These index sets induce respective partitions on the kernel matrix K and the label vector y (we shall use the lower-case letters s, e, o and r for such partitions).

By writing out the Kuhn-Tucker conditions (2)–(3) before and after an update $\Delta\alpha$ we obtain the following condition that must be satisfied after an update:

$$\begin{bmatrix} \Delta g_c \\ \Delta g_s \\ \Delta g_r \\ 0 \end{bmatrix} = \begin{bmatrix} y_c & K_{cs} \\ y_s & K_{ss} \\ y_r & K_{rs} \\ 0 & y_s^T \end{bmatrix} \underbrace{\begin{bmatrix} \Delta\mu \\ \Delta\alpha_s \end{bmatrix}}_{\Delta s} + \Delta\alpha_c \begin{bmatrix} K_{cc}^T \\ K_{cs}^T \\ K_{cr}^T \\ y_c \end{bmatrix}. \tag{4}$$

One can see that $\Delta\alpha_c$ is in equilibrium with $\Delta\alpha_s$ and μ: any change to $\Delta\alpha_c$ must be absorbed by the appropriate changes in $\Delta\alpha_s$ and μ in order for the condition (4) to hold.

The main equilibrium condition (4) can be further refined as follows. It follows from (2) that $\Delta g_s = 0$. Then lines 2 and 4 of the system (4) can be re-written as

$$\begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 & \alpha_s^T \\ \alpha_s & K_{ss} \end{bmatrix} \Delta s + \begin{bmatrix} \alpha_c \\ K_{cs}^T \end{bmatrix} \Delta\alpha_c. \tag{5}$$

This linear system is easily solved for Δs as follows:

$$\Delta s = \beta \Delta\alpha_c, \tag{6}$$

where

$$\beta = -\underbrace{\begin{bmatrix} 0 & \alpha_s^T \\ \alpha_s & K_{ss} \end{bmatrix}^{-1}}_{Q} \underbrace{\begin{bmatrix} \alpha_c \\ K_{cs}^T \end{bmatrix}}_{\vec{\eta}} \tag{7}$$

is the gradient of the manifold of optimal solutions parameterized by α_c.

One can further substitute (6) into lines 1 and 3 of the system (4):

$$\begin{bmatrix} \Delta g_c \\ \Delta g_r \end{bmatrix} = \gamma \Delta \alpha_c, \tag{8}$$

where

$$\gamma = \begin{bmatrix} y_c & K_{cs} \\ y_r & K_{rs} \end{bmatrix} \beta + \begin{bmatrix} K_{cc} \\ K_{cr}^T \end{bmatrix} \tag{9}$$

is the gradient of the manifold of gradients g_r at an optimal solution parameterized by α_c.

The upshot of these derivations is that the update is controlled by very simple sensitivity relations (6) and (8), where β is sensitivity of Δs with respect to $\Delta \alpha_c$ and γ is sensitivity of $\Delta g_{c,r}$ with respect to $\Delta \alpha_c$.

2.3 Accounting

Unfortunately the system (4) cannot be used directly to obtain the new SVM state. The problem lies in the changing composition of the sets S and R with the change of Δs and $\Delta \alpha_c$ in Eq. (4). To handle this problem, the main strategy of the algorithm is to identify the largest increase $\Delta \alpha_c$ such that some point migrates between the sets S and R. Four cases must be considered to account for such structural changes:

1. Some α_i in S reaches a bound (an upper or a lower one). Let ε be a small number. Compute the sets

$$I_+^S = \{i \in S : \beta_i > \varepsilon\}$$
$$I_-^S = \{i \in S : \beta_i < -\varepsilon\}.$$

The examples in set I_+^S have positive sensitivity with respect to the weight of the current example; that is, their weight would increase by taking the step $\Delta \alpha_c$.[2] These examples should be tested for reaching the upper bound C. Likewise, the examples in set I_-^S should be tested for reaching zero. The examples with $-\varepsilon < \beta_i < \varepsilon$ should be ignored, as they are insensitive to $\Delta \alpha_c$. Thus the possible weight updates are

$$\Delta \alpha_i^{\max} = \begin{cases} C - \alpha_i, & \text{if } i \in I_+^S \\ -\alpha_i, & \text{if } i \in I_-^S, \end{cases}$$

and the largest possible $\Delta \alpha_c^S$ before some example in S moves to R is

$$\Delta \alpha_c^S = \operatorname*{absmin}_{i \in I_+^S \cup I_-^S} \frac{\Delta \alpha_i^{\max}}{\beta_i}, \tag{10}$$

where

$$\operatorname{absmin}(x) := \min_i |x_i| \cdot \operatorname{sign}(x_{(\operatorname{argmin}_i |x_i|)}).$$

2. Some g_i in R reaches zero. Compute the sets

$$I_+^R = \{i \in E : \gamma_i > \varepsilon\}$$
$$I_-^R = \{i \in O : \gamma_i < -\varepsilon\}.$$

---

2. It can be shown that the step $\Delta \alpha_c$ is always positive in the incremental case.

The examples in set I_+^R have positive sensitivity of the gradient with respect to the weight of the current example. Therefore their (negative) gradients can potentially reach zero. Likewise, gradients of the examples in set I_-^R are positive but are pushed towards zero with the increasing weight of the current example. Thus the largest increase $\Delta\alpha_c^g$ before some point in R moves to S can be computed as

$$\Delta\alpha_c^R = \min_{i \in I_+^R \cup I_-^R} \frac{-g_i}{\gamma_i}. \tag{11}$$

3. g_c becomes zero. This case is similar to case 2, with the feasibility test in the form

$$\gamma_c > \varepsilon.$$

If the update is feasible the largest step $\Delta\alpha_c^g$ is computed as

$$\Delta\alpha_c^g = \frac{-g_c}{\gamma_c}. \tag{12}$$

4. α_c reaches C. The largest possible increment $\Delta\alpha_c^\alpha$ is clearly

$$\Delta\alpha_c^\alpha = C - \alpha_c. \tag{13}$$

Finally, the smallest of the four values

$$\Delta\alpha_c^{\max} = \min\left(\Delta\alpha_c^S, \Delta\alpha_c^R, \Delta\alpha_c^g, \Delta\alpha_c^\alpha\right) \tag{14}$$

constitutes the largest possible increment of α_c.

After the largest possible increment of α_c is determined, the updates Δs and Δg are carried out. The inverse matrix Q must be also re-computed in order to account for the new composition of the set S. Efficient update of this matrix is presented in section 2.4. The sensitivity vectors β and γ must be also re-computed. The process repeats until the gradient of the current example becomes zero or its weight reaches C. The high-level summary of incremental SVM algorithm is given in Algorithm 1.

2.4 Recursive Update of the Inverse Matrix

It is clearly infeasible to explicitly invert the matrix Q in Eq. (7) every time the set S is changed. Luckily, it is possible to use the fact that this set is always updated one element at a time – an example is either added to, or removed from the set S.

Consider addition first. When the example x_k[3] is added to the set S, the matrix to be inverted is partitioned as follows:

$$R := \begin{bmatrix} 0 & y_s^T & y_k \\ y_s & K_{ss} & K_{ks}^T \\ y_k & K_{ks} & K_{kk} \end{bmatrix} = \begin{bmatrix} Q^{-1} & \eta_k \\ \eta_k^T & K_{kk} \end{bmatrix}, \tag{15}$$

where

$$\eta_k = \begin{bmatrix} y_k \\ K_{ks}^T \end{bmatrix}.$$

---

3. A different index k is used to emphasize that this example is not the same as the current example x_c.

Algorithm 1 Incremental SVM algorithm: high-level summary.

1: Read example x_c, compute g_c.
2: **while** $g_c < 0$ & $\alpha_c < C$ **do**
3:   Compute β and γ according to (7), (9).
4:   Compute $\Delta\alpha_c^S, \Delta\alpha_c^R, \Delta\alpha_c^g$ and $\Delta\alpha_c^\alpha$ according to (10)–(13).
5:   Compute $\Delta\alpha_c^{\max}$ according to (14).
6:   $\alpha_c \leftarrow \alpha_c + \Delta\alpha_c^{\max}$.
7:   $\alpha_s \leftarrow \beta\Delta\alpha_c^{\max}$
8:   $g_{c,r} \leftarrow \gamma\Delta\alpha_c^{\max}$
9:   Let k be the index of the example yielding the minimum in (14).
10:   **if** $k \in S$ **then**
11:     Move k from S to either E or O.
12:   **else if** $k \in E \cup O$ **then**
13:     Move k from either E or O to S.
14:   **else**
15:     $\{k = c$: do nothing, the algorithm terminates.$\}$
16:   **end if**
17:   Update Q recursively. $\{$See section 2.4.$\}$
18: **end while**

Define $\beta_k = -Q\eta_k$, and denote the enlarged inverse matrix by $\tilde{Q}$. Applying the Sherman-Morrison-Woodbury formula for block matrix inversion (see e.g. Golub and van Loan (1996)) to $\tilde{Q}$, we obtain:

$$
\underbrace{\begin{bmatrix} Q^{-1} & \eta_k \\ \eta_k^T & K_{kk} \end{bmatrix}}_{\tilde{Q}}^{-1} = \begin{bmatrix} Q + \kappa^{-1}(Q\eta_k)(Q\eta_k)^T & -\kappa^{-1}Q\eta_k \\ -\kappa^{-1}(Q\eta_k)^T & \kappa^{-1} \end{bmatrix}
$$

$$
= \begin{bmatrix} Q + \kappa^{-1}\beta_k\beta_k^T & \kappa^{-1}\beta_k \\ \kappa^{-1}\beta_k^T & \kappa^{-1} \end{bmatrix}
$$

$$
= \begin{bmatrix} Q & 0 \\ 0 & 0 \end{bmatrix} + \frac{1}{\kappa}\begin{bmatrix} \beta_k \\ 1 \end{bmatrix}\begin{bmatrix} \beta_k^T & 1 \end{bmatrix}, \tag{16}
$$

where

$$
\kappa = K_{kk} - \eta_k^T Q\eta_k. \tag{17}
$$

Thus the update of the inverse matrix involves expansion with a zero row and column and addition of a rank-one matrix obtained via a matrix-vector multiplication. The running time needed for an update of the inverse matrix is quadratic in the size of Q, which is much better than explicit inversion.

The removal case is straightforward. Knowing the inverse matrix $\tilde{Q}$ and using (16), we can write

$$
\tilde{Q} = \begin{bmatrix} q_{11} & q_{12} \\ q_{21} & q_{22} \end{bmatrix} = \begin{bmatrix} Q + \kappa^{-1}\beta_k\beta_k^T & \kappa^{-1}\beta_k \\ \kappa^{-1}\beta_k^T & \kappa^{-1} \end{bmatrix}.
$$

Re-writing this block matrix expression component-wise, we have:

$$Q = q_{11} - \kappa^{-1}\beta_k\beta_k^T \tag{18}$$

$$\beta_k = \kappa q_{12} = q_{12}/q_{22} \tag{19}$$

$$\beta_k^T = \kappa q_{21} = q_{21}/q_{22} \tag{20}$$

$$\kappa^{-1} = q_{22}. \tag{21}$$

Substituting the last three relations into the first we obtain:

$$Q = q_{11} - \frac{q_{12}q_{21}}{q_{22}}. \tag{22}$$

The running time of the removal operations is also quadratic in the size of Q.

3. Extension to One-Class Classification

Availability of labels, especially online, may not be possible in certain applications. For example, analysis of security logs is extremely time-consuming, and labels may be available only in limited quantities after forensic investigation. As another example, monitoring of critical infrastructures, such as power lines or nuclear reactors, must prevent a system from reaching a failure state which may lead to gravest consequences. Nevertheless, algorithms similar to SVM classification can be applied for data-description, i.e. for automatic inference of a concept descriptions deviations from which are to be considered abnormal. Since examples of the class to be learned are not (or rarely) available, the problem is known as "one-class classification".

The two well-known approaches to one-class classification are separation of data points from the origin (Schölkopf et al., 2001) and spanning of data points with a sphere (Tax and Duin, 1999). Although they use different geometric constructions, these approaches lead to similar, and in certain cases even identical, formulations of dual optimization problems. As we shall see, incremental learning can be naturally extended to both of these approaches.

The dual formulation of the "sphere" one-class classification is given by the following quadratic program:

$$\max_{\alpha} \quad \sum_{i=1}^{n}\alpha_i k(x_i,x_i) - \frac{1}{2}\sum_{i=1}^{n}\sum_{j=1}^{n}\alpha_i\alpha_j k(x_i,x_j)$$

$$\text{subject to:} \quad 0 \leq \alpha_i \leq C,\ i=1,\ldots,n \tag{23}$$

$$\sum_{i=1}^{n}\alpha_i = 1.$$

In order to extend the C&P algorithm for this problem consider the following abstract saddle-point quadratic problem:

$$\max_{\mu}\ \min_{\substack{0\leq\alpha\leq C \\ a^T\alpha+b=0}} W := -c^T\alpha + \frac{C}{2}\alpha^T H\alpha + \mu(a^T\alpha+b). \tag{24}$$

It can be easily seen that formulation (24) generalizes both problems (23) and (1), subject to following definition of the abstract parameters:

$$c = \text{diag}(K), \quad H = K, \quad a = 1, \quad b = 1.$$

Algorithm 2 Initialization of incremental one-class SVM.

1: Take the first $\lfloor \frac{1}{C} \rfloor$ objects, assign them weight C and put them in E.
2: Take the next object c, assign $\alpha_c = 1 - \lfloor \frac{1}{C} \rfloor C$ and put it in S.
3: Compute the gradients g_i of all objects, using Eq. (2).
4: Compute μ so as to ensure non-positive gradients in E: $\mu = -\max_{i \in E} g_i$
5: Enter the main loop of the incremental algorithm.

The C&P algorithm presented in sections 2.2–2.4 can be applied to formulation (24) almost without modification. The only necessary adjustment is a special initialization procedure for identification of an initial feasible solution presented in Algorithm 2.[4]

4. Convergence Analysis

A very attractive feature of Algorithm 1 is that it obviously makes progress to an optimal solution if a non-zero update $\Delta\alpha_c$ is found. Thus potential convergence problems arise only when a zero update step is encountered. This can happen in several situations. First, if the set S is empty, no non-zero update of $\Delta\alpha_c$ is possible since otherwise the equality constraint of the SVM training problem is violated. Second, a zero update can occur when two or more points simultaneously migrate between the index sets. In this case, each subsequent point requires a structural update without improvement of $\Delta\alpha_c$. Furthermore, it must be guaranteed, that after a point migrates from one set to another, say from E to S, it is not immediately thrown out after the structure and the sensitivity parameters are re-computed. These issues constitute the scope of convergence analysis to be covered in this section. In particular we present a technique for handling the special case of the empty set S and show that immediate cycling is impossible if a kernel matrix is positive semi-definite.

4.1 Empty Set S

The procedure presented in the previous two sections requires a non-empty set of unbounded support vectors, otherwise a zero matrix must be inverted in Eq. (7). To handle this situation, observe that the main instrument needed to derive the sensitivity relations is pegging of the gradient to zero for unbounded support vectors, which follows from the Kuhn-Ticker conditions (2). Notice that the gradient can also be zero for some points in sets E and O. Therefore, if the set S is empty we can freely move some examples with zero gradients from the sets E and O to it and continue from there. The question arises: what if no points with zero gradients can be found in E or O?

With $\Delta\alpha = 0$, $\Delta\alpha_c = 0$ and no examples in the set S, the equilibrium condition (4) reduces to

$$\Delta g_c = y_c \Delta\mu$$
$$\Delta g_r = y_r \Delta\mu. \tag{25}$$

This is an equilibrium relation between $\Delta g_{c,r}$ and the scalar $\Delta\mu$, in which sensitivity is given by the vector $[y_c; y_r]$. In other words, one can change μ freely until one of components in g_r or g_c hits zero, which would allow an example to be brought into S.

---

4. Through the presentation of the C&P algorithm it was assumed that a feasible solution is always available. This is indeed no problem for the classification SVM since the zero solution is always feasible. For the sphere formulation of one-class classification this is not the case.

The problem remains – since $\Delta\mu$ is free as opposed to non-negative $\Delta\alpha_c$ – to determine the direction in which the components of g_r are pushed by changes in μ. This can be done by first solving (25) for $\Delta\mu$, which yields the dependence of Δg_r on Δg_c:

$$\Delta g_r = -\frac{y_r}{y_c}\Delta g_c.$$

Since Δg_c must be non-negative (gradient of the current example is negative and should be brought to zero if possible), the direction of Δg_r is given by $-\frac{y_r}{y_c}$. Hence the feasibility conditions can be formulated as

$$I_+^R = \{i \in E : -\frac{y_i}{y_c} > \varepsilon\}$$
$$I_-^R = \{i \in O : -\frac{y_i}{y_c} < -\varepsilon\}.$$

The rest of the argument essentially follows the main case. The largest possible step $\Delta\mu^R$ is computed as

$$\Delta\mu^R = \min_{i \in I_+^R \cup I_-^R} \frac{-g_i}{y_i}. \tag{26}$$

and the largest possible step $\Delta\mu^c$ is computed as

$$\Delta\mu^c = -\frac{g_c}{y_c}. \tag{27}$$

Finally, the update $\Delta\mu^{\max}$ is chosen as

$$\Delta\mu^{\max} = \min\left(\Delta\mu^R, \Delta\mu^c\right). \tag{28}$$

4.2 Immediate Cycling

A more dangerous situation can occur if an example entering the set S is immediately thrown out at a next iteration without any progress. It is not obvious why this kind of immediate cycling cannot take place, since the sensitivity information contained in vectors β and γ does not seem to provide a clue what would happen to an example after the structural change. The analysis provided in this section shows that, in fact, such look-ahead information is present in the algorithm and that this property is related to positive semi-definiteness of a kernel matrix.

When an example is added to the set S, the corresponding entry in the sensitivity vector β is added at the end of this vector. Let us compute $\tilde{\beta}$ *after an addition of example k to set S*:

$$\tilde{\beta} = -\tilde{Q}\eta = -\begin{bmatrix} Q & 0 \\ 0 & 0 \end{bmatrix}\begin{bmatrix} y_c \\ K_{cs} \end{bmatrix} - \frac{1}{\kappa}\begin{bmatrix} \beta_k\beta_k^T & \beta_k \\ \beta_k^T & 1 \end{bmatrix}\begin{bmatrix} y_c \\ K_{cs} \end{bmatrix}.$$

Computing the last line in the above matrix products we obtain:

$$\tilde{\beta}_{\text{end}} = -\frac{1}{\kappa}\Bigg(\underbrace{\beta_k^T\begin{bmatrix} y_c \\ K_{cs\backslash\text{end}} \end{bmatrix} + K_{ck}}_{\gamma_k}\Bigg). \tag{29}$$

Let us assume that κ (defined in Eq. (17)) is non-negative. Examples with $\kappa = 0$ must be prevented from entering the set S externally, otherwise invertibility of Q is ruined; therefore $\kappa > 0$

for the examples entering the set S. Thus we can see that the sign of $\tilde{\beta}_{end}$ after the addition of the element k to set S is *the opposite* to the sign of γ_k before the addition. Recalling the feasibility conditions for inclusion of examples in set S, one can see that if an example is joining from the set O its $\tilde{\beta}$ after inclusion will be positive, and if an example is joining from the set E its $\tilde{\beta}$ after inclusion will be negative. Therefore, in no case will an example be immediately thrown out of the set S.

Thus to prove that immediate cycling is impossible it has to be shown that $\kappa \geq 0$. Two technical lemmas are useful before we proceed with the proof of this fact.

Lemma 1 *Let $z = -Q\eta_k$, $\tilde{z} = [z,1]^T$. Then $\kappa = \tilde{z}^T R \tilde{z}$.*

Proof By writing out the quadratic form we obtain:

$$\tilde{z}^T R \tilde{z} = z^T Q^{-1} z + \eta_k^T z + z^T \eta_k + K_{kk}$$
$$= \eta_k^T Q Q^{-1} Q \eta_k - 2\eta_k Q \eta_k + K_{kk}$$
$$= K_{kk} - \eta_k Q \eta_k.$$

The result follows by the definition of κ. ∎

The next lemma establishes a sufficient condition for non-negativity of a quadratic form with a matrix of the special structure possessed by matrix R.

Lemma 2 *Let $\tilde{x} = [x_0, x]^T$, $\tilde{K} = \begin{bmatrix} 0 & y^T \\ y & K \end{bmatrix}$, where x_0 is a scalar, x, y are vectors of length n, and K is a positive semi-definite $n \times n$ matrix. If $x^T y = 0$ then $\tilde{x}^T \tilde{K} \tilde{x} \geq 0$.*

Proof By writing out the quadratic form we obtain

$$\tilde{x}^T \tilde{K} \tilde{x} = 2x_0 x^T y + x^T K x.$$

Since K is positive semi-definite the second term is greater than or equal to 0, whereas the first term vanished by the assumption of the lemma. ∎

Finally, the intermediate results of the lemmas are used in the main theorem of this section.

Theorem 3 *If the kernel matrix K is positive semi-definite then $\kappa \geq 0$.*

Proof Lemma 1 provides a quadratic form representation of κ. Our goal is thus to establish its non-negativity using the result of Lemma 2.

Using the partition matrix inversion formula we can write Q (cf. Eq. 7) as

$$Q = \begin{bmatrix} -\frac{1}{\delta} & \frac{1}{\delta} y_s^T K_{ss}^{-1} \\ \frac{1}{\delta} K_{ss}^{-1} y_s & K_{ss}^{-1} - \frac{1}{\delta} K_{ss}^{-1} y_s y_s^T K_{ss}^{-1} \end{bmatrix},$$

where $\delta = y_s^T K_{ss}^{-1} y_s$. Substituting Q into the definition of z in Lemma 1 and explicitly writing out the first term, we obtain:

$$z := \begin{bmatrix} z_0 \\ z_{\backslash 0} \end{bmatrix} = - \begin{bmatrix} -\frac{1}{\delta} & \frac{1}{\delta} y_s^T K_{ss}^{-1} \\ \frac{1}{\delta} K_{ss}^{-1} y_s & K_{ss}^{-1} - \frac{1}{\delta} K_{ss}^{-1} y_s y_s^T K_{ss}^{-1} \end{bmatrix} \begin{bmatrix} y_k \\ K_{sk} \end{bmatrix}$$

$$= \begin{bmatrix} \frac{1}{\delta} y_k - \frac{1}{\delta} y_s^T K_{ss}^{-1} K_{sk} \\ -\frac{1}{\delta} K_{ss}^{-1} y_s y_k - K_{ss}^{-1} K_{sk} + \frac{1}{\delta} K_{ss}^{-1} y_s y_s^T K_{ss}^{-1} K_{sk} \end{bmatrix}.$$

Then

$$z_{\backslash 0}^T y_s + y_k = -\frac{1}{\delta} y_k \underbrace{y_s^T K_{ss}^{-1} y_s}_{\delta} - K_{sk}^T K_{ss}^{-1} y_s + \frac{1}{\delta} K_{sk}^T K_{ss}^{-1} y_s \underbrace{y_s^T K_{ss}^{-1} y_s}_{\delta} + y_k = 0$$

and the result follows by Lemma 2. ∎

5. Runtime Analysis and Efficient Design

Let us now zoom in on computational complexity of Algorithm 1 which constitutes a minor iteration of the overall training algorithm.[5] Asymptotically, the complexity of a minor iteration is quadratic in a number of examples learned so far: re-computation of the gradient, β and γ involve matrix-vector multiplications, which have quadratic complexity, and the recursive update of an inverse matrix has also been shown (cf. Section 2.4) to be quadratic in the number of examples. These estimates have to be multiplied by a number of minor iterations needed to learn an example. The number of minor iterations depends on the structure of a problem, namely on how often examples migrate between the index sets until an optimal solution is found. This number cannot be control in the algorithm, and can be potentially exponentially large.[6]

For the practical purposes it is important to understand the constants hidden in asymptotic estimates. To this end, the analysis of a direct implementation of Algorithm 1 in Matlab is presented in Section 5.1. The focus of our analysis lies on the complexity of the main steps of a minor iteration in terms of arithmetic and memory access operations. This kind of analysis is important because arithmetics can be implemented much more efficiently than memory access. Performance-tuned numeric libraries, such as BLAS[7] or ATLAS,[8] make extensive use of the cache memory which is an order of magnitude faster than the main memory. Therefore, a key to efficiency of the incremental SVM algorithm lies in identifying performance bottlenecks associated with memory access operations and trying to eliminate them in a clever design. The results of our analysis are illustrated by profiling experiments in Section 5.2, in which relative complexity of the main computational operations, as a percentage of total running time, is measured for different kernels and sample sizes.

---

5. A major iteration corresponds to inclusion of a new example; therefore, all complexity estimates must be multiplied by a number of examples to be learned. This is, however, a (very) worst case scenario, since no minor iteration is needed for many points that do not become support vectors at the time of their inclusion.

6. The structure of the problem is determined by the geometry of a set of feasible solutions in a feature space. Since we essentially follow the outer boundary of the set of feasible solutions, we are bound by the same limitation as linear and non-linear programming in general, for which it is known that problems exist with exponentially many vertices in a set of feasible solutions (Klee and Minty, 1972).

7. http://www.netlib.org/blas/

8. http://math-atlas.sourceforge.net/

It follows from our analysis and experiments that the main difficulty in efficient implementation of incremental SVM indeed lies in selection of non-contiguous elements of matrices, e.g. in Eq. (9). The problem cannot be addressed within Matlab in which storage organization is one-dimensional. Furthermore, it must be realized, as our experience showed us, that merely re-implementing incremental SVM without addressing the tradeoff between selection and arithmetic operations, for example using C++ with one-dimensional storage, *does not* solve the problem. The solution proposed in Section 5.3, which allows to completely eliminate expensive selection operations at a cost of minor increase of arithmetic operations, is based on a mixture of row- and column-major storage and on the gaxpy-type (e.g. Golub and van Loan (1996)) matrix-vector products. The evaluation of the new design presented in Section 5.4 shows performance improvement of 5 to 20 times.

5.1 Computational Complexity of Incremental SVM

On the basis of the pseudo-code of Algorithm 1 we will now discuss the key issues that will later be used for a more efficient implementation of the incremental SVM.

Line 1: Computation of g_c is done according to Eq. (2). This calculation requires partial computation of the kernel row K_{cs} for the current example and examples in the set S. If the condition of the while loop in line 2 does not hold then the rest of the kernel row K_{cr} has to be computed for Eq. (9) in line 3. Computation of a kernel row is expensive since a subset of input points, usually stored as a matrix, has to be selected using the index sets S and R.

Line 3: The computation of γ via Eq. (9) is especially costly since a two-dimensional selection has to be performed to obtain the matrix K_{rs} ($O(sr)$ memory access operations), followed by a matrix-vector multiplication ($O(sr)$ arithmetic operations). The computation of β in line 3 is relatively easy because the inverse matrix Q is present and only the matrix-vector multiplication for Eq. (7) ($O(ss)$ arithmetic operations) has to be performed. The influence of γ and β for the algorithm scales with the size of set S and the number of data points.

Lines 4-8: These lines have minor runtime relevance because only vector-scalar calculations and selections are to be performed.

Lines 9-16: Administration operations for sets S and R have inferior complexity. If a kernel row of the example k is not present (in case of x_k entering the S from O) then it has to be re-computed for the update of the inverse matrix in line 17.

Line 17: The update of the inverse matrix requires the calculation of κ according to Eq. (17) ($O(ss)$ arithmetic operations), which is of a similar order as the computation of β. The expansion and rank-one matrix computation have also effect on the algorithm runtime. The expansion requires memory operation (in the naive implementation, a reallocation of the entire matrix) and is thus expensive for a large inverse matrix.

To summarize, the main performance bottlenecks of the algorithm are lines 1, 3, and 17, in which memory access operations take place.

5.2 Performance Evaluation

We now proceed with experimental evaluation of the findings of Section 5.1. As a test-bed the MNIST handwritten digits data set[9] is used to profile the training of an incremental SVM. For every digit, a test run is made on the data sets of size 1000, 2000, 3000, 5000 and 10,000 randomly drawn from the training data set. Every test run was performed for a linear kernel, a polynomial kernel of degree 2 and an RBF kernel with $\sigma = 30$. Profiles were created by the Matlab profiler.

Eight operations were identified where a Matlab implementation of Algorithm 1 spends the bulk of its runtime (varying from 75% to 95% depending on a digit). Seven of these operations pertain to the bottlenecks identified in Section 5.1. Another relatively expensive operation is augmentation of a kernel matrix with a kernel row. Figure 1 shows proportions of runtime spent in the most expensive operations for the digit 8.
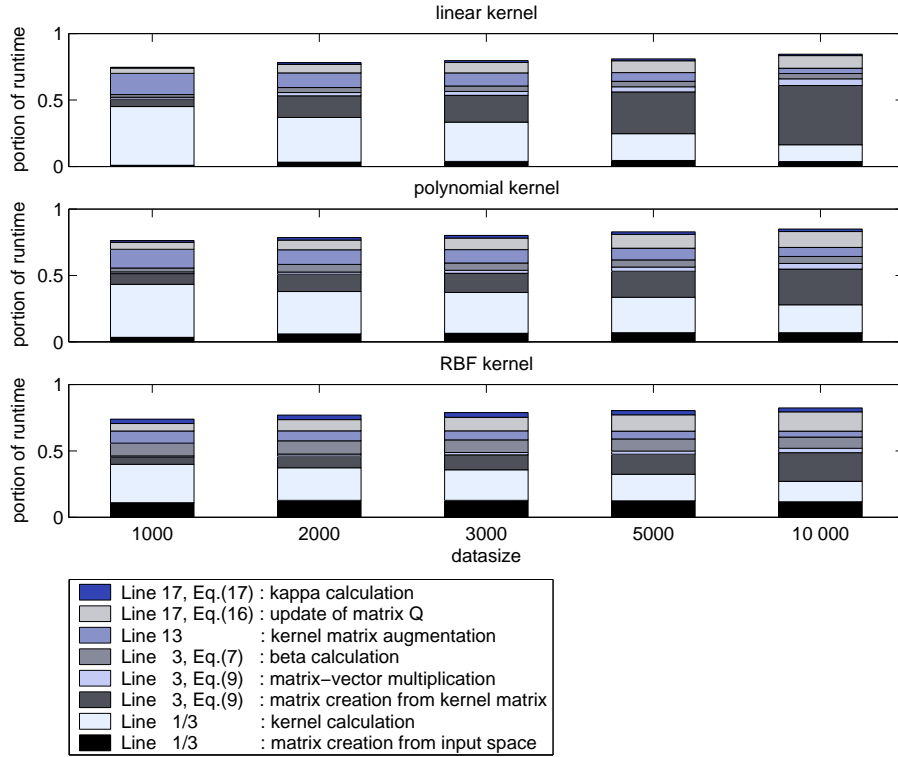
Figure 1: Profiling results for digit 8 (MNIST).

The analysis clearly shows that memory access operations dominate the runtime shown in Figure 1. It also reveals that the portion of kernel computation scales down with increasing data size for all three kernels.

---

9. This data set can be found at http://yann.lecun.com/exdb/mnist/, and contains 60,000 training and 10,000 test images of size 28×28.

5.3 Organization of Matrix Storage and Arithmetic Computations

Having found the weak spots in the Matlab implementation of the incremental SVM, we will now consider the possibilities for the efficiency improvement. In to order gain control over the storage design we choose C++ as an implementation platform. Numerical operations can be efficiently implemented by using the ATLAS library.

5.3.1 STORAGE DESIGN

As it was mentioned before, the difficulty of selection operations in Matlab result from storing matrices as one-dimensional arrays. For this type of storage, selection of rows and columns necessarily required a large amount of copying.

Figure 2: C++ matrix design

An alternative representation of a matrix in C++ is a pointer-to-pointer scheme shown in Figure 2. Depending on whether rows or columns of a matrix are stored in one-dimensional pointed to `double**` pointers, either a row-major or a column-major storage is realized.

What are the benefits of a pointer-to-pointer storage for our purposes? Such matrix representation has the advantage that selection along the pointer-of-pointer array requires hardly any time since only addresses of rows or columns need to be fetched. As a result one can completely eliminate selection operations in line 1 by storing the input data in a column-major matrix. Furthermore, by storing a kernel matrix in a row-major format (a) additions of kernel rows can be carried out without memory relocation, and (b) selections in the matrix K_{rs} are somewhat optimized since $r \gg s$. Another possible problem arises during addition of a new example when columns have to added to a kernel matrix. This problem can be solved by pre-allocation of columns which can be done in constant-time (amortized).

5.3.2 MATRIX OPERATIONS

The proposed memory design does not completely solve the problem: we still have to perform column selection in s when computing K_{rs}. Although tolerable for smaller number of support vectors,

---

Algorithm 3 C++ calculation of γ for (9) using gaxpy-type matrix-vector multiplication

---

Create an empty vector z
for $i = 1 : |s|$ **do**
  $z = \beta_{i+1} K_{s_i,:} + z$ {gaxpy call}
end for
Compute $\gamma = \beta_1 y_r + z_r + K_{cr}$

Algorithm 4 C++ calculation of γ for (9) using a naive matrix-vector multiplication

---

Create a matrix Z
$Z = \left[y_r^T ; K_{sr} \right]$
Compute $\gamma = \beta^T Z + K_{cr}$ {dgemm call}

---

the problem becomes acute when s grows with the arrival of more and more examples. Yet it turns out to be possible to eliminate even this selection by re-organizing the computation of γ.

Consider the following form of computing Eq. (9):[10]

$$\gamma^T = \beta_1 y_r^T + \beta_{2:\mathrm{end}}^T K_{sr} + K_{cr}. \tag{30}$$

The first and the last terms are merely vectors, while the middle term is computed using a matrix-vector multiplication and selection over a matrix. Using the transposed matrix K_{sr} (still stored in a row-major form) at first seems counter-intuitive, as we argued in the previous section that expensive selection should be carried out over short indices in S and not the long indices in R. However, consider the following observation:

> Since $s \ll r$ we can just as well run the product $\beta_{2:\mathrm{end}}^T K_{s,1:\mathrm{end}}$ at a tolerable extra cost of $O(ss)$ arithmetic operations.

By doing so we do not need to worry about selection! The extra s elements in a product vector can be discarded (of course by a selection, however selection over a vector is cheap).

Still another problem with the form (30) of the γ-update remains. If we run it as an inner-product update, i.e. multiplying a row vector with columns of a matrix stored in a row-major format, this loop must be run over the elements in non-contiguous memory. This is as slow as using selection. However, by running it as a gaxpy-type update (e.g. Golub and van Loan, 1996) we end up with loops running over the *rows* of a kernel matrix, which brings a full benefit of performance-tuned numerics. The summary of the gaxpy-type computation of γ is given in Algorithm 3. For comparison, Algorithm 4 shows a naive implementation of the γ-update using inner-product operations (dgemm-type update).

This same construction can be also applied to the inverse matrix update in Eq. (16) by using the following format:

$$Q_i = \frac{\beta_i}{\kappa} \beta + Q_{i,:}. \tag{31}$$

By doing so explicit creation of a rank-one matrix can be avoided.

To summarize our design, by using the row-major storage of (transposed) matrix K and the gaxpy-type matrix-vector multiplication selection can be avoided by at a cost of extra $O(ss)$ arithmetic operations and performing a selection on a resulting vector.

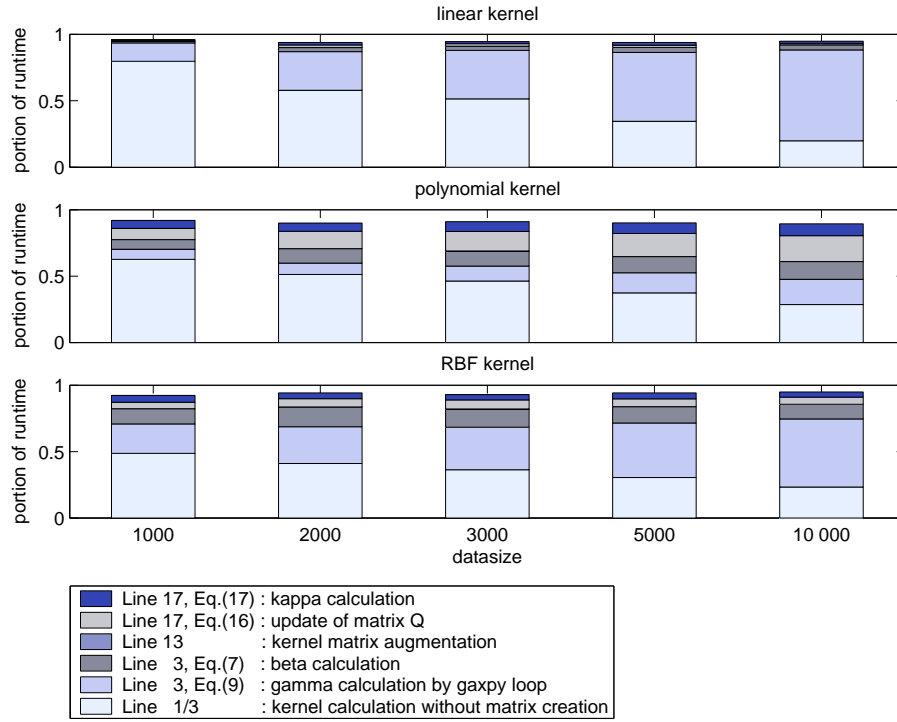10. For cleaner notation we ignore the first line in Eq. (9) here.

Figure 3: C++ runtime proportion digit 8 (MNIST)

5.4 Experimental Evaluation of the New Design

The main goal of the experiments to be presented in this section is to evaluate the impact of the new design of storage and arithmetic operations on the overall performance of incremental SVM learning. In particular, the following issues are to be investigated:

- How is the runtime profile of the main operations affected by the new design?

- How does the overall runtime of the new design scale with an increasing size of a training set?

- What is the overall runtime improvement over the previous implementations and how does it depend on the size of a training set?

To investigate the runtime profiles of the new design, check-pointing has been realized in our C++ implementation. The profiling experiment presented in Section 5.2 (cf. Figure 1) has been repeated for the new design, and the results are shown in Figure 3. The following effects of the new design can be observed:

1. The selection operation in Lines 1/3 is eliminated. The selection was necessary in a Matlab implementation due to one-dimensional storage – a temporary matrix had to be created in order to represent a sub-matrix of the data. In the new design using the column-major storage for the data matrix a column sub-matrix can be directly passed (as pointers to columns) without a need for selection.

2. Matrix creation has been likewise eliminated in the computation of γ in Line 3. However, the relative cost of the gaxpy computation in the new design remains as high as the relative cost of the combined matrix creation / matrix-vector multiplication operations in the Matlab implementation. The relative cost of the β computation is not affected by the new design.

3. Matrix augmentation in Line 13 takes place at virtually no computational cost – due to row-major storage of the kernel matrix.

4. The relative cost of operations in Line 17 is not affected by the new design.

The overall cost distribution of main operations remains largely the same in the new design, kernel computation having the largest weight for small training sets and gamma computation – for the large training sets. However, as we will see from the following experiments, the new design results in major improvement of the absolute running time.

Evaluation of the absolute running time is carried out by means of the scaling factor experiments. The same data set and the same SVM parameters are used as in the profiling experiments. Four implementations of incremental SVM are compared: the original Matlab implementation of C&P (with leave-one-out error estimation turned off), the Matlab implementation of Algorithm 1, the C++ implementation of Algorithms 1 & 3 and the C++ implementation of Algorithms 1 & 4. The latter configuration is used in order to verify that the performance gains indeed stem from the gaxpy-type updates rather than from switching from Matlab to C++.

The algorithms are run on the data sets ranging from 1000 to 10000 examples in size, and the training times are plotted against the training set size at a log-log scale. These plots are shown in Figure 4 (for the linear kernel) and 5 (for the RBF kernel). The results for the polynomial kernel are similar to the linear kernel and are not shown. Ten plots are shown separately for each of the digits.

One can see that the C++ implementation significantly outperforms both Matlab implementations. The RBF kernel is more difficult for training than the linear kernel for the MNIST data set, which is reflected by a larger proportion of support vectors (on average 15% for the RBF kernel compared to 5% with the linear kernel, at 10000 training examples). Because of this the experiments with the C&P algorithm at 10000 training points were aborted. The Matlab implementation of Algorithm 1 was able to crank about 15000 examples with the RBF kernel, whereas the C++ implementation succeeded to learn 28000 examples, before running out of memory for storing the kernel matrix and the auxiliary data structures (at about 3GB). The relative performance gain of the C++ implementation using gaxpy-updates against the "best Matlab competitor" and against the dgemm-updates is shown in Figure 6 (linear kernel, C&P algorithm) and Figure 7 (RBF kernel, Algorithm 1). Major performance improvement in comparison to Matlab and the naive C++ implementations can be observed, especially visible on larger training set sizes.

6. Applications

As it was mentioned in the introduction, various applications of incremental SVM learning can be foreseen. Two exemplare applications are presented in this session in order to illustrate some potential application domains.
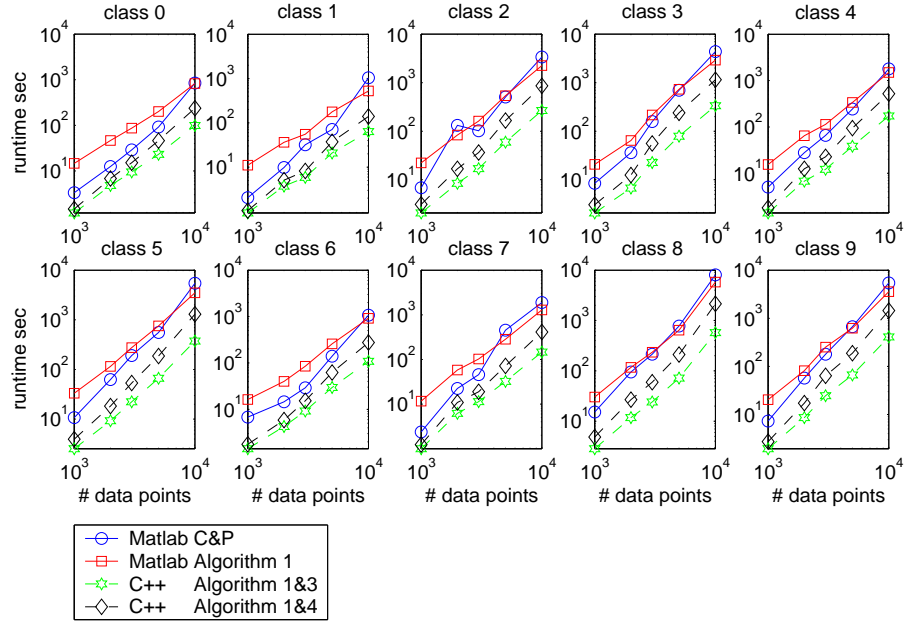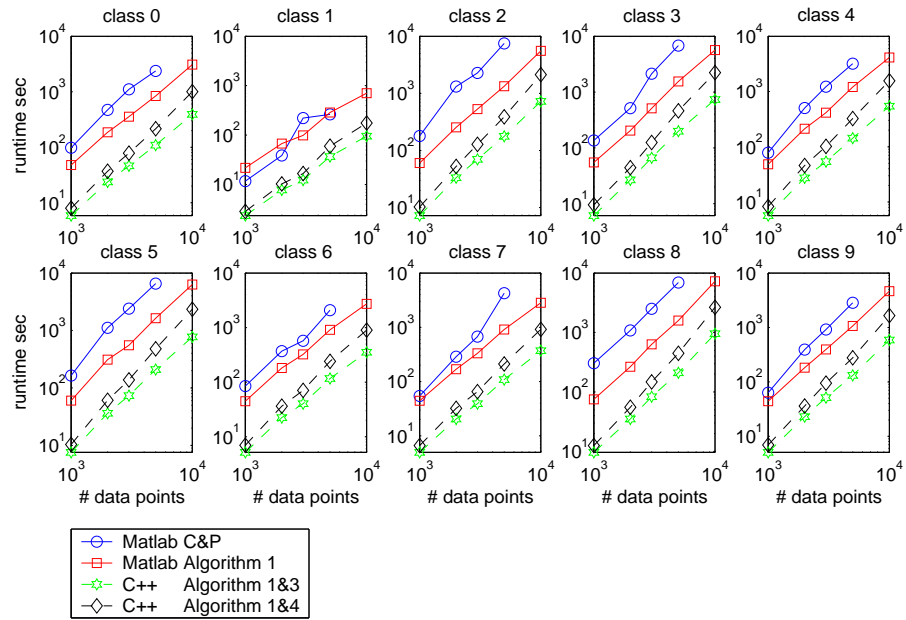
Figure 4: Scaling factor plots of incremental SVM algorithms with the linear kernel.

Figure 5: Scaling factor plots of incremental SVM algorithms with the RBF kernel ($\sigma = 30$).
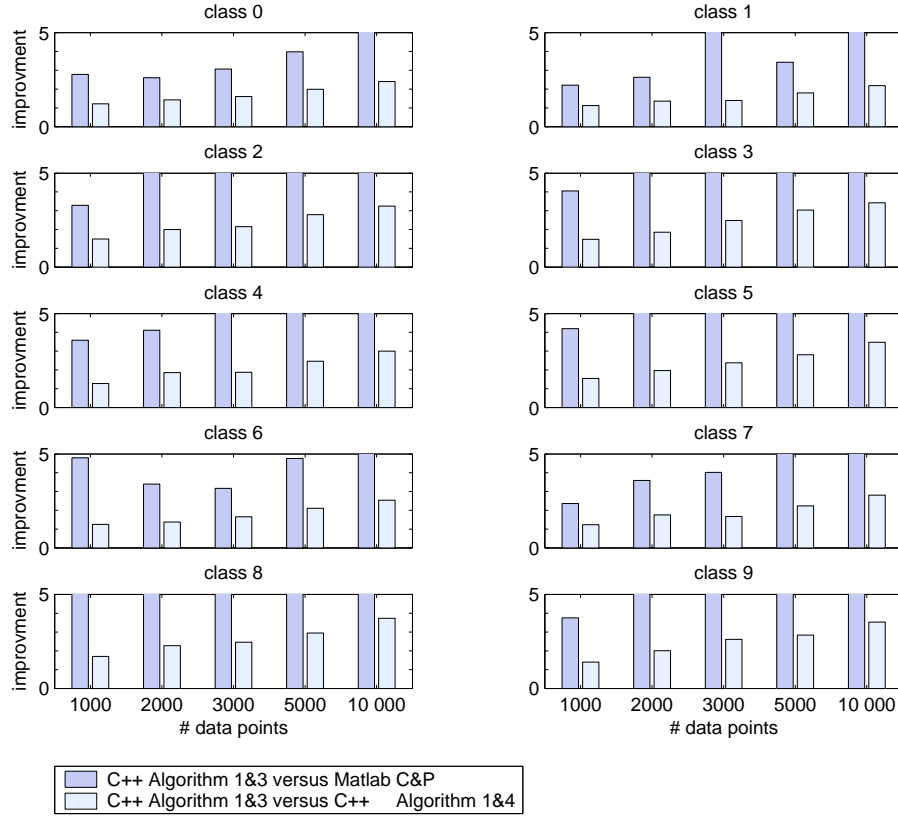
Figure 6: Runtime improvement, linear kernel.

6.1 Learning with Limited Resources

To make SVM learning applicable to very large data sets, a classifier has to be constrained to have a limited number of objects in memory. This is, in principle, exactly what an online classifier with fixed window size *M* does. Upon arrival of a new example, a least relevant example needs to be removed before a new example can be incorporated. A reasonable criterion for relevance is the value of the weight.

USPS experiment: learning with limited resources. As a proof of concept for learning with limited resources we train an SVM on the USPS data set under the limitations on the number of points that can be seen at a time. The USPS data set contains 7291 training and 2007 images of handwritten digits, size 16×16 (Vapnik, 1998). On this 10-class data set 10 support vector classifiers with a RBF kernel, $\sigma^2 = 0.3 \cdot 256$ and $C = 100$, were trained.[11] During the evaluation of a new object, it is assigned to the class corresponding to the classifier with the largest output. The total classification error on the test set for different window sizes *M* is shown in Figure 8.

One can see that the classification accuracy deteriorates marginally (by about 10%) until the working size of 150, which is about 2% of the data. True, by discarding "irrelevant" examples, one removes potential support vectors that cannot be recovered at a later stage. Therefore one can

---

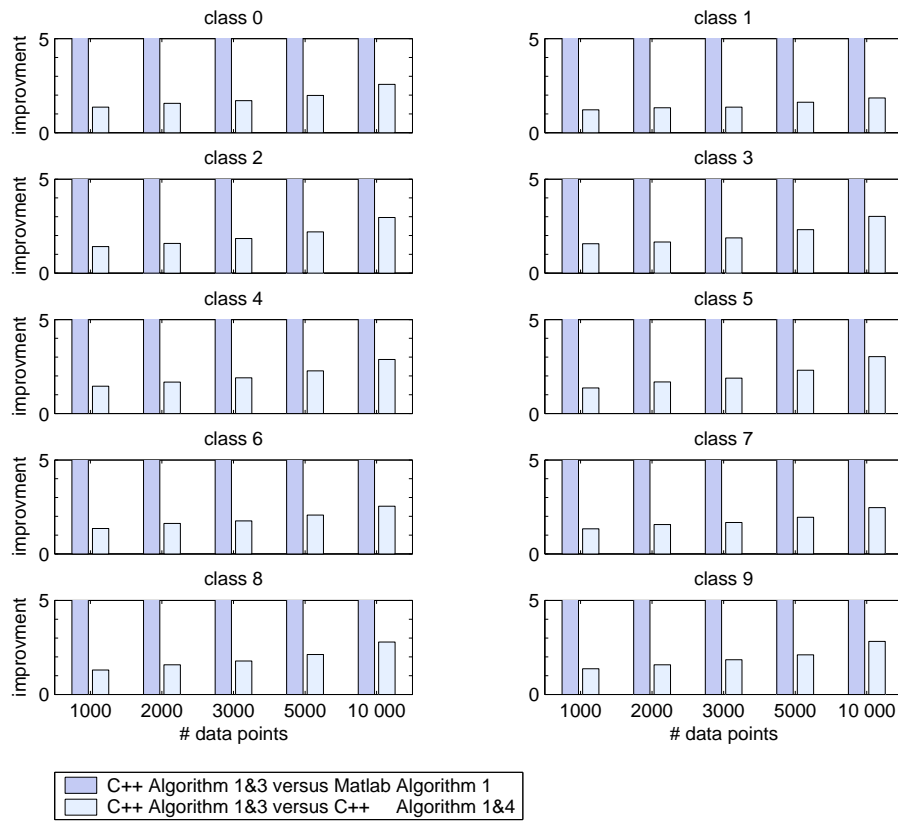11. The best model parameters as reported in (Vapnik, 1998) were used.

Figure 7: Runtime improvement, RBF kernel.

Figure 8: Test classification errors on the USPS data set, using a support vector classifier (RBF kernel, $\sigma^2 = 0.3 \cdot 256$) with a limited "window" of training examples.

expect that performance of a limited memory classifier would be worse than that of an unrestricted classifier. It is also obvious that no more points than the number of support vectors are eventually needed, although the latter number is not known in advance. The average number of support vectors

per each unrestricted 2-class classifier in this experiment is 274. Therefore the results above can be interpreted as reducing the storage requirement by 46% from the minimal at the cost of 10% increase of classification problem.

Notice that the proposed strategy differs from the caching strategy, typical for many SVM[light]-like algorithms (Joachims, 1999; Laskov, 2002; Collobert and Bengio, 2001), in which kernel products are re-computed if the examples are found missing in the fixed-size cache and the accuracy of the classifier is not sacrificed. Our approach constitutes a trade-off between accuracy and computational load because kernel products never need to be re-computed. It should be noted, however, that computational cost of re-computing the kernels can be very significant, especially for the problems with complicated kernels such as string matching or convolution kernels.

6.2 Active Learning

Another promising application of incremental SVM is active learning. In this scenario, instead of having all data labelled beforehand, an algorithm "actively" chooses examples for which labels must be assigned by a user. Active learning can be extremely successful, if not indispensable, when labelling is expensive, e.g. in computer security or in drug discovery applications.

A very powerful active learning algorithm using SVM was proposed by Warmuth et al. (2003). Assume that the goal of learning is to identify "positive" examples in a data set. The meaning of positivity can vary across applications; for example, it can be binding properties of molecules in drug discovery applications, or hacker attacks in security applications. Selection of a next point to be labelled is carried out in the algorithm of Warmuth et al. (2003) using two heuristics that can be derived from an SVM classifier trained on points with known labels. The "largest positive" heuristic selects the point that has the largest classification score among all examples still unlabeled. The "near boundary" heuristic selects the point whose classification score has the smallest absolute value. Although the semantics of these two heuristics differ – in one case we trying to explore the space of positive examples as fast as possible, whereas in the other case the effort is focused on learning the boundary – in both cases the SVM has to be re-trained after each selection. In the original application of Warmuth et al. (2003) the data samples were relatively small, therefore one could afford re-training SVM from scratch after addition of new points. Obviously, a better way to proceed is by applying incremental learning as presented in this paper.

In the remaining part of this section experiments will be presented that prove the usefulness of active learning in the intrusion detection context. Since the observed data can contain thousands and even millions of examples it is clear that the problem can be addressed only using incremental learning. As a by-product of our experiments, it will be seen that active learning helps to uncover the structure of a learning problem revealed by the number of support vectors.

The underlying data for the experiments is taken from the KDD Cup 1999 data set.[12] As a training set, 1000 examples are randomly drawn with an attack rate of 10 percent. The incremental SVM was run with the linear kernel and the RBF kernel with $\sigma = 30$. An independent set of the same length with the same attack distribution is used for testing. The results are not averaged over multiple repetitions in order not to "disturb" the semantics of different phases of active learning as can be seen from the ROC curves. However, similar behavior was observed over multiple experiments.

KDD Cup experiment: active learning.   Consider the following learning scenario. Assume that we can run an anomaly detection tool over our data set which ranks all the points according to

---

12. http://www-cse.ucsd.edu/users/elkan/clresults.html
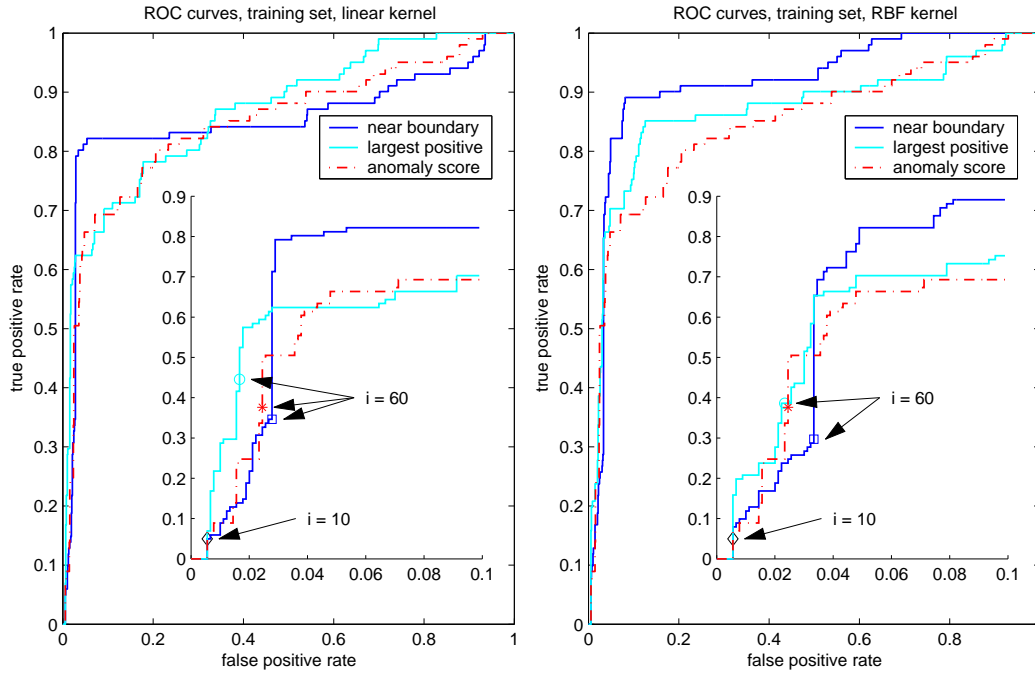
Figure 9: linear kernel, n=10 and m=50

their degree of anomaly. No labels are needed for this; however, if we know them, we can evaluate anomaly detection by a ROC curve. It is now the goal to use active learning to see if a ROC curve of anomaly detection can be improved.

We take the first n examples with the highest anomaly scores and train a (batch) SVM to obtain an initial model. After that we turn to active learning and learn the next m examples. We are now ready to classify the remaining examples by the trained SVM. The question arises: after spending manual effort to label $n+m$ examples, can we classify the remaining examples better than anomaly detection?

In order to address this question, an accuracy measure must be defined for our learning scenario. This can be done using the fact that only the ranking of examples according to their scores – and not the score values themselves – matters for the computation of a ROC curve (Cortes and Mohri, 2004). The ranking in our experiment can be defined as follows: the first n examples are ranked according to their anomaly scores, the next m examples are ranked according to their order of inclusion during the active learning phase, and the remaining examples are ranked according to their classification scores.

The ROC curves for active learning with the two heuristics and for anomaly detection are shown in Figure 9. One can easily see a different behavior exhibited by the two active learning rules. The "largest positive" rule attains the highest true positive rate during the active learning phase, but does not perform significantly better than anomaly detection during the classification phase ($i > 60$). On the contrary, the "near boundary" rule is close or worse than anomaly detection during the learning phase but exhibit a sharp increase of the true positive rate after moving to classification mode. Its accuracy then remains consistently better than anomaly detection for a considerable false positive interval (until FP = 0.3 for the linear kernel and until FP = 0.9 for the RBF kernel). Similar behavior

Figure 10: linear kernel, n=10 and m=50

of the two heuristics in the active learning phase was also observed by Warmuth et al. (2003). Yet the "near boundary" heuristic is obviously more suitable for classification, since it explores the boundary region and not merely the region of positive examples.

Another interesting insight can be gained by investigating the behavior of active learning on test data. In this case, supervised learning can also be drawn into comparison. In particular, we consider a full SVM (using a training set size of 1000 examples as opposed to only 60 examples used in the active learning) and a reduced SVM. The latter is obtained from a full SVM by finding a hyperplane closest to a full SVM hyperplane subject to 1-norm regularization over expansion coefficients (cf. Schölkopf et al. (1999)). The regularization constant is chosen such that a reduced SVM has approximately the same number of support vectors as the solution obtain by active learning (in our case the value $\lambda = 2.5$ resulted in about 30 support vectors). Thus one can compare active learning with supervised learning given equal complexity of solutions.

The ROC curves of active learning, supervised learning and anomaly detection on test data are shown in Figure 10. It can be observed that the "near-boundary" heuristic of active learning attains a solution which is at least as good (for FP ≤ 0.4) as a reduced SVM for the linear kernel and significantly better for the RBF kernel. This shows that active learning does a very good job at discovering the necessary structure of a solution – it picks a better representation within a desired complexity since it is using a learning-related criterion to select an interesting representation instead of a merely geometric one.

7. Discussion and Conclusions

Online learning algorithms have proved to be essential when dealing with (a) very large (see e.g. Le-Cun et al. (1998); Bordes et al. (2005); Tsang et al. (2005)) or (b) non-stationary data (see e.g. Robbins and Munro (1951); Murata (1992); Murata et al. (1997, 2002)). While classical neural networks (e.g. LeCun et al. (1998); Saad (1998)) have a well established online learning toolbox for optimization, incremental learning techniques for Support Vector Machines have been only recently developed (Cauwenberghs and Poggio, 2001; Tax and Laskov, 2003; Martin, 2002; Ma et al., 2003; Ma and Perkins, 2003).

The current paper contributes two-fold to the field of incremental SVM learning. The convergence analysis of the algorithm has been performed showing that immediate cycling of the algorithm is impossible provided a kernel matrix is positive semi-definite. Furthermore, we propose a better scheme for organization of memory and arithmetic operations in exact incremental SVM using the gaxpy-type updates of the sensitivity vector. As it is demonstrated by our experiments, the new design results in major constant improvement in the running time of the algorithm.

The achieved performance gains open wide possibilities for application of incremental SVM to various practical problems. We have presented exemplary applications to two possible scenarios: learning with limited resources and active learning. Potential applications of incremental SVM learning include, among others, drug discovery, intrusion detection, network surveillance, monitoring of non-stationary time series etc. Our implementation is available free of charge for academic use at http://www.mind-ids.org/Software.

It is interesting to compare exact incremental learning to recently proposed alternative approaches to online learning. The recent work of Bordes et al. (2005) presents an online algorithm for L1 SVM, in which a very close approximation of the exact solution is built online before the last gap is bridged in the REPROCESS phase in an offline fashion. This algorithm has been shown to scale well to several hundred thousand examples, however its online solution is not as accurate as the exact solution. It has been observed (cf. Fig. 9 in Bordes et al. (2005)) that the REPROCESS phase may result in major improvement of the test error and may come at a high price in comparison with the online phase, depending on a data set. Another recent algorithm, the Core Vector Machine of Tsang et al. (2005), is based on the L2 formulation of an SVM and has be shown to scale to several million of examples. The idea of this algorithm is to approximate a solution to an L2 SVM by a solution to the two-class Maximal Enclosing Ball problem, for which several efficient online algorithms are known. While scalability results of CVM are very impressive, the approximation of the exact solution can likewise in higher test errors.

The major limilation of the exact incremental learning is its memory requirement, since the set of support vectors must be retained in memory during the entire learning. Due to this limitation, the algorithm is unlikely to be scalable beyond tens of thousands examples; however, for data sizes within this limit it offers an advantage of immediate availablility of the exact solution (crucial in e.g learning of non-stationary problems) and reversibility.

Future work will include further investigation of properties of incremental SVM such as numerical stability and their utility for tracking the values of generalization bounds. A relationship with parametric optimization needs to be further clarified. Extensions to advance learning modes, such as learning for structured domains and semi-supervised learning, are being considered.

Acknowledgments

References

D. Angluin. Queries and concept learning. *Machine Learning*, 2:319–342, 1988.

C. M. Bishop. *Neural Networks for Pattern Recognition*. Oxford University Press, 1995.

B.. Blankertz, G. Dornhege, C. Schäfer, R. Krepki, J. Kohlmorgen, K.-R. Müller, V. Kunzmann, F. Losch, and G. Curio. BCI bit rates and error detection for fast-pace motor commands based on single-trial EEG analysis. *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, 11:127–131, 2003.

A. Bordes, S. Ertekin, J. Wesdon, and L. Bottou. Fast kernel classifiers for online and active learning. *Journal of Machine Learning Research*, 6:1579–1619, 2005.

G. Cauwenberghs and T. Poggio. Incremental and decremental support vector machine learning. In T. K. Leen, T. G. Dietterich, and V. Tresp, editors, *Advances in Neural Information Processing Systems*, volume 13, pages 409–415. MIT Press, 2001.

S. Chakrabarti. *Mining the Web: Discovering Knowledge from Hypertext Data*. Morgan-Kaufmann, 2002. ISBN 1-55860-754-4.

C.-C. Chang and C.-J. Lin. Libsvm: Introduction and benchmarks. Technical report, Department of Computer Science and Information Engineering, National Taiwan University, Taipei, 2000.

R. Collobert and S. Bengio. SVMTorch: Support vector machines for large-scale regression problems. *Journal of Machine Learning Research*, 1:143–160, 2001.

C. Cortes and M. Mohri. AUC optimization vs. error rate minimization. In *Proc. NIPS'2003*, 2004.

E. Eskin, A. Arnold, M. Prerau, L. Portnoy, and S. Stolfo. *Applications of Data Mining in Computer Security*, chapter A geometric framework for unsupervised anomaly detection: detecting intrusions in unlabeled data. Kluwer, 2002.

G. H. Golub and C. F. van Loan. *Matrix Computations*. John Hopkins University Press, Baltimore, London, 3rd edition, 1996.

T. Joachims. Making large–scale SVM learning practical. In B. Schölkopf, C. J. C. Burges, and A. J. Smola, editors, *Advances in Kernel Methods — Support Vector Learning*, pages 169–184, Cambridge, MA, 1999. MIT Press.

J. Kivinen, A. J. Smola, and R. C. Williamson. Online learning with kernels. In T. G. Diettrich, S. Becker, and Z. Ghahramani, editors, *Advances in Neural Inf. Proc. Systems (NIPS 01)*, pages 785–792, 2001.

F. Klee and G. J. Minty. How good is the simplex algorithm? In O. Sisha, editor, *Inequalities III*, pages 159–175. Academic Press, 1972.

P. Laskov. Feasible direction decomposition algorithms for training support vector machines. *Machine Learning*, 46:315–349, 2002.

P. Laskov, C. Schäfer, and I. Kotenko. Intrusion detection in unlabeled data with quarter-sphere support vector machines. In *Proc. DIMVA*, pages 71–82, 2004.

Y. LeCun, L. Bottou, G. B. Orr, and K.-R. Müller. Efficient backprop. In G. Orr and K.-R. Müller, editors, *Neural Networks: Tricks of the Trade*, volume 1524, pages 9–53, Heidelberg, New York, 1998. Springer LNCS.

N. Littlestone, P. M. Long, and M. K. Warmuth. On-line learning of linear functions. Technical Report CRL-91-29, University of California at Santa Cruz, October 1991.

J. Ma and S. Perkins. Time-series novelty detection using one-class Support Vector Machines. In *IJCNN*, 2003. to appear.

J. Ma, J. Theiler, and S. Perkins. Accurate online support vector regression. http://nis-www.lanl.gov/˜jt/Papers/aosvr.pdf, 2003.

M. Martin. On-line Support Vector Machines for function approximation. Technical report, Universitat Politècnica de Catalunya, Departament de Llengatges i Sistemes Informàtics, 2002.

N. Murata. *A statistical study on the asymptotic theory of learning*. PhD thesis, University of Tokyo (In Japanese), 1992.

N. Murata, M. Kawanabe, A. Ziehe, K.-R. Müller, and S.-I. Amari. On-line learning in changing environments with applications in supervised and unsupervised learning. *Neural Networks*, 15 (4-6):743–760, 2002.

N. Murata, K.-R. Müller, A. Ziehe, and S. i. Amari. Adaptive on-line learning in changing environments. In M. C. Mozer, M. I. Jordan, and T. Petsche, editors, *Advances in Neural Information Processing Systems*, volume 9, page 599. The MIT Press, 1997.

G. Orr and K.-R. Müller, editors. *Neural Networks: Tricks of the Trade*, volume 1524. Springer LNCS, 1998.

J. Platt. Fast training of support vector machines using sequential minimal optimization. In B. Schölkopf, C. J. C. Burges, and A. J. Smola, editors, *Advances in Kernel Methods — Support Vector Learning*, pages 185–208, Cambridge, MA, 1999. MIT Press.

L. Ralaivola and F. d'Alché Buc. Incremental support vector machine learning: A local approach. *Lecture Notes in Computer Science*, 2130:322–329, 2001.

H. Robbins and S. Munro. A stochastic approximation method. *Ann. Math. Stat.*, 22:400–407, 1951.

S. Rüping. Incremental learning with support vector machines. Technical Report TR-18, Universität Dortmund, SFB475, 2002.

D. Saad, editor. *On-line learning in neural networks*. Cambridge University Press, 1998.

B. Schölkopf, S. Mika, C. J. C. Burges, P. Knirsch, K.-R. Müller, G. Rätsch, and A. J. Smola. Input space vs. feature space in kernel-based methods. *IEEE Transactions on Neural Networks*, 10(5): 1000–1017, September 1999.

B. Schölkopf, J. Platt, J. Shawe-Taylor, A. J. Smola, and R. C. Williamson. Estimating the support of a high-dimensional distribution. *Neural Computation*, 13(7):1443–1471, 2001.

N. A. Syed, H. Liu, and K. K. Sung. Incremental learning with support vector machines. In *SVM workshop, IJCAI*, 1999.

D. Tax and R. Duin. Data domain description by support vectors. In M. Verleysen, editor, *Proc. ESANN*, pages 251–256, Brussels, 1999. D. Facto Press.

D. M. J. Tax and P. Laskov. Online SVM learning: from classification to data description and back. In C. et al. Molina, editor, *Proc. NNSP*, pages 499–508, 2003.

I. Tsang, J. Kwok, and P.-M. Cheung. Core Vector Machines: fast SVM training on very large data sets. *Journal of Machine Learning Research*, 6:363–392, 2005.

V. N. Vapnik. *Statistical Learning Theory*. Wiley, New York, 1998.

M. K. Warmuth, J. Liao, G. Rätsch, M. Mathieson, S. Putta, and C. Lemmem. Support Vector Machines for active learning in the drug discovery process. *Journal of Chemical Information Sciences*, 43(2):667–673, 2003.